

**Subject:** ML slides pack - proximity measures by Claudio Sartori

**Course:** Artificial Intelligence - LM

**Department:** DISI (Department of Computer Science and Engineering) - University of Bologna, Italy

**Author:** Roberto Zanolli

## Proximity measures

Proximity measures are numerical measures of how alike, different or distant two data objects are. This is fundamental in many applications for ML, in particular clustering.

### Similarity

Numerical measure of how alike two data objects are.

- is higher when objects are more alike.
- often falls in the range [0,1]

### Dissimilarity

Numerical measure of how different are two data objects.

- lower when objects are more alike
- minimum dissimilarity is often 0
- upper limit varies

### Proximity

Refers to a similarity or dissimilarity.

- almost same as distance

### Common definitions by attribute type:

Attribute type	Dissimilarity	Similarity
Nominal ↳ No sequence = No distance	$d = \begin{cases} 0 & \text{if } p = q \\ 1 & \text{if } p \neq q \end{cases}$	$s = \begin{cases} 1 & \text{if } p = q \\ 0 & \text{if } p \neq q \end{cases}$
Ordinal Values mapped to integers 0 to V-1	$d = \frac{ p-q }{V-1}$	$s = 1 - \frac{ p-q }{V-1}$
Interval or Ratio	$d =  p - q $	$s = \frac{1}{1+d}$ or $s = 1 - \frac{d - \min(d)}{\max(d) - \min(d)}$

## Euclidean distance – L2

$$dist = \sqrt{\sum_{d=1}^D (p_d - q_d)^2}$$

Where  $D$  is the number of dimensions (attributes) and  $p_d$  and  $q_d$  are, respectively, the  $d$ -th attributes (components) of data objects  $p$  and  $q$

**Important:** Standardization/rescaling is necessary if scales differ

## Minkowski distance – L $r$

$$dist = \left( \sum_{d=1}^D |p_d - q_d|^r \right)^{\frac{1}{r}}$$

Where  $D$  is the number of dimensions (attributes) and  $p_d$  and  $q_d$  are, respectively, the  $d$ -th attributes (components) of data objects  $p$  and  $q$

**Important:** Standardization/rescaling is necessary if scales differ

$r$  is a parameter which is chosen depending on the data set and the application

Common choices for  $r$ :

$r = 1$ : also named city block, **Manhattan**,  $L_1$  norm

- it is the best way to discriminate between zero distance and near zero distance
- a 1 change on any coordinate causes a 1 change in the distance (unlike Euclidean where the effect is diluted by the square root of the sum of squares)
- works better than Euclidean in very high dimensional spaces

$r = 2$ : **Euclidean**,  $L_2$  norm

$r = \infty$ : also named **Chebyshev**, supremum,  $L_{max}$  norm,  $L_\infty$  norm

- considers only the dimension where the difference is maximum
- provides a simplified evaluation, disregarding the dimensions with lower differences

$$dist_\infty = \lim_{r \rightarrow \infty} \left( \sum_{d=1}^D |p_d - q_d|^r \right)^{\frac{1}{r}} = \max_d |p_d - q_d|$$

## Mahalanobis distance

Considers **data distribution**

The Mahalanobis distance between two points  $p$  and  $q$  decreases if, *keeping the same Euclidean distance*, the segment connecting the points is stretched along a direction of greater variation of data.

The distribution is described by the covariance matrix of the data set

$$dist_m = \sqrt{(p - q) \Sigma^{-1} (p - q)^T}$$

$$\Sigma_{ij} = \frac{1}{N-1} \sum_{k=1}^N (x_{ki} - \bar{x}_i)(x_{kj} - \bar{x}_j)$$

## Mahalanobis distance – example

$$\Sigma = \begin{bmatrix} 0.3 & 0.2 \\ 0.2 & 0.3 \end{bmatrix}$$

$$A = (0.5, 0.5)$$

$$B = (0, 1)$$

$$C = (1, 1)$$

The Euclidean distances AB and AC are the same

$$dist_m(A, B) = 2.236068$$

$$dist_m(A, C) = 1$$

## Covariance matrix

Variation of pairs of random variables

- the summation is over all the observations
- the main diagonal contains the variances
- the values are positive if the two variables grow together
- if the matrix is diagonal the variables are non-correlated
- if the variables are standardised the diagonal contains "one"
- if the variables are standardised and non correlated, the matrix is the identity and the Mahalanobis distance is the same as the Euclidean

## Common properties of a distance

### 1. Positive definiteness:

$$Dist(p, q) \geq 0 \quad \forall p, q$$

and

$\text{Dist}(p, q) = 0$  if and only if  $p = q$

2. **Symmetry:**  $\text{Dist}(p, q) = \text{Dist}(q, p)$

3. **Triangle inequality:**  $\text{Dist}(p, q) \leq \text{Dist}(p, r) + \text{Dist}(r, q)$ ,  $\forall p, q, r$

A distance function satisfying all the properties above is called a **metric**

### EXAMPLE OF INVALID METRIC:

SITUATION: I have 2 events, one is today at 4 PM while the other is at 3 PM of tomorrow.

- the distance is not  $|6 - 4| = |3 - 4| = 1$  h
- the distance is 23 h

↳ this distance is an invalid metric  
for this use case

## Recommendation system example: movie ratings

Why Manhattan works better:

- sparse matrix: users rate only a few movies
- better handling of missing values
- more robust to extreme ratings

## Distance calculation example

Comparing user A and user D (common movies: Inception, Matrix):

Manhattan distance:

$$|5 - 5| + |4 - 3| = 0 + 1 = 1$$

Euclidean distance:

$$\sqrt{(5 - 5)^2 + (4 - 3)^2} = \sqrt{0 + 1} = 1$$

For vectors [5,1,5,1] and [4,2,4,2]:

$$\text{Manhattan: } |5 - 4| + |1 - 2| + |5 - 4| + |1 - 2| = 4$$

$$\text{Euclidean: } \sqrt{1^2 + 1^2 + 1^2 + 1^2} = 2$$

Manhattan better captures total distance between preferences

## Practical application: collaborative filtering

Algorithm:

- Calculate Manhattan distance from user A to all other users
- Identify the k nearest neighbors (e.g., k=3)
- Recommend to A the movies appreciated by neighbors that A hasn't seen yet

### Example:

If user D is most similar to user A

Recommend "Avengers" to A (D rated it 4 stars)

This approach works particularly well with Manhattan distance for sparse datasets

---

## Supremum (Chebyshev) distance use cases

- **Anomaly detection:** efficient for detecting extreme deviations across high-dimensional features
  - **Chessboard distance:** used in modeling real-world systems where maximum single-step moves matter
  - **Industrial applications:** applied in quality control, where maximum tolerances are checked
- 

## Cosine similarity use cases

- **Text mining and document similarity:** used for document comparison and recommendation systems
  - **Image similarity:** applied in image retrieval systems to match images with similar features
  - **Recommendation systems:** collaborative filtering methods leverage cosine similarity
- 

## Manhattan vs cosine distance: key differences

Manhattan distance	Cosine similarity
Measures absolute differences	Measures angle/direction
Sensitive to magnitude	Magnitude-invariant
Formula: $\sum_{i=1}^n  x_i - y_i $	Formula: $\frac{x \cdot y}{\ x\  \ y\ }$
Range: $[0, \infty]$	Range: $[-1, 1]$

---

## Example 1: movie ratings (when magnitude matters)

Three users rate movies on scale 1-5:

User	Movie 1	Movie 2	Movie 3	Profile
Alice	5	5	5	Loves everything
Bob	3	3	3	Moderate ratings
Carol	5	5	4	Loves most things

### Manhattan distance:

- Alice-Bob:  $|5 - 3| + |5 - 3| + |5 - 3| = 6$
- Alice-Carol:  $|5 - 5| + |5 - 5| + |5 - 4| = 1$
- Bob-Carol:  $|3 - 5| + |3 - 5| + |3 - 4| = 5$

**Result:** Alice is closest to Carol (both are enthusiastic raters)

### Cosine similarity perspective:

All users are highly similar! Cosine ignores that Bob rates lower overall.

### Use Manhattan when rating scale/magnitude matters

---

## Example 2: document similarity (when direction matters)

Three documents represented by word frequencies:

Document	"machine"	"learning"	"algorithm"	"data"
Doc A	10	8	12	15
Doc B	2	1	3	2
Doc C	5	12	3	8

### Manhattan distance:

- A-B:  $|10 - 2| + |8 - 1| + |12 - 3| + |15 - 2| = 37$
- A-C:  $|10 - 5| + |8 - 12| + |12 - 3| + |15 - 8| = 25$

### Cosine similarity:

- A-B: 0.959
- A-C: 0.785

**Result:** Cosine correctly identifies Doc B as more similar to A (same topic despite length difference)

### Use Cosine when you care about proportional similarity, not absolute counts

---

## Example 3: user behavior vectors

Two users' activity on a website (clicks per section):

User	News	Sports	Tech	Entertainment
User X	100	80	60	40
User Y	10	8	6	4

**Manhattan distance:** 270 (very large - treats them as different)

**Cosine similarity:** 1.0 (perfect - same interest distribution!)

### Decision:

- use Manhattan if engagement level matters
- use Cosine if interest pattern matters

---

## Quick decision guide

Scenario	Manhattan	Cosine
Absolute values matter	✓	
Scale/magnitude important	✓	
Different-length documents		✓
Interest patterns		✓

Rule of thumb:

- Manhattan: "How different are the values?"
- Cosine: "How similar are the patterns?"

## Jaccard similarity

**Definition:** measures overlap between two sets relative to their union

Use cases:

- document similarity in information retrieval
- image processing
- clustering and community detection